



Samueli
Computer Science



Uncertainty based Reinforcement Learning with Function Approximation

Weitong Zhang

Department of Computer Science, UCLA

Apr 18, 2022

Overview of my research

- Area: Reinforcement learning and online learning
- Reinforcement Learning with function approximation
 - Representation learning in reinforcement learning [ZHZZG21]
 - Exploration with function approximation [ZZG21, NeurIPS] (this talk)
 - Convergence guarantee of popular RL algorithm [WZXG20, NeurIPS]
- Deep learning based recommendation systems [ZZLG20, ICLR], [JZZGW21, ICLR]
- Machine Learning in interdisciplinary fields
 - COVID-19 forecasting [CRLB21, PNAS]
 - Mechanistic deciphering for Chemistry reactions

Reinforcement learning is useful

- Example: self-driving car
- Each time you will:
 - Get information from sensors, cameras, GPS...
 - Take actions: brake, accelerate or make turns
- Goal: reach the designation safely and efficiently



Image Credit: scharfsinn86 - stock.adobe.com

Reinforcement learning is interesting

- Example: chess game
- Each time you will:
 - Get information on board
 - Take actions:
 - Be2? Bd2? h3?
- Goal: Defend your king and checkmate your opponent




King's Indian Defense, known for Black's aggressive attack, was a big challenge for White until RL / search algorithms emerged
Screen shot from chess.com

Reinforcement learning: formal definition

- Markov Decision Processes (MDPs)
- For each time step $h = 1, 2, \dots, H$, the agent will:
 - Observe state s_h
 - Take action $a_h = \pi_h(s_h)$ by policy π
 - Receive reward $r_h(s_h, a_h)$
 - Transit to next state $s_{h+1} \sim \mathbb{P}(\cdot | s_h, a_h)$
- Goal: maximize the cumulative rewards:

$$Q_h^\pi(s, a) = \mathbb{E} \left[\sum_{h'=h}^H r_h(s, a) | s, a \right], V_h^\pi(s) = Q_h^\pi(s, \pi(s))$$



By Megajuice - Own work, CC0, <https://commons.wikimedia.org/w/index.php?curid=57895741>

Towards large state space using function approximation


- Tabular methods are infeasible in practice
- Go game has approximately 10^{360} states
- Deep neural networks can perfectly extract features of the states
- RL with function approximation widely studied [JYWJ19, YW19, ZGS20]




By Goban1 - Own work, Public Domain,
<https://commons.wikimedia.org/w/index.php?curid=15223468>

Uncertainty in RL with function approximation

- Uncertainty is omnipresent!
 - Errors in the model...
 - Noise in the data...
 - Missing information...
- Uncertainty is important!
 - Performance issue
 - Efficiency issue
 - Fairness issue...



Uncertainty from an inaccurate model (Using Taylor approx.)




Uncertainty from the noise in the model

ID	Color	Weight	Broken	Class
1	Black	80	Yes	1
2	Yellow	100	No	2
3	Yellow	120	Yes	2
4	Blue	90	No	2
5	Blue	85	No	2
6	?	60	No	1
7	Yellow	100	?	2
8	?	40	?	1

Uncertainty from missing data

$$\begin{array}{ll} \text{1% better every day} & 1.01^{365} = 37.18 \\ \text{1% worse every day} & 0.99^{365} = 0.03 \end{array}$$



Minor uncertainty on beginning can lead to significant performance issue

Tasks in RL with Function Approximation


- Exploration: Pure exploration without reward signals (this talk)
- Model Misspecification Issue: Control the approximation error in the model
- Representation Learning: Select good representation to improve performance
- Partially observed RL and non-Markovian RL: Missing information in current observation
- Fairness in RL: Make fair decision when uncertainty exists
- Deep RL: quantify the uncertainty in neural networks used in RL

Key technical issue: How to precisely quantify and utilize the uncertainty in RL with function approximation?

Using uncertainty to guide exploration in RL

Efficiency Challenge in Reinforcement Learning


- Learning the environment (model) can be inefficient!
- Challenge: Can we reuse the model on different goals?
- E.g. a map (**same environment**) can be used to navigate to different destination (**different targets**)



We can use the same map of LA to get to DTLA, LAX or Hollywood from UCLA!

Why conventional reward-driven RL fails?

- Reward-guided exploration cannot well explore environment!
- Reward signal might not be given during the exploration.
 - *You will never know where you will go when constructing the map!*



Reward-guided exploration with target LAX cannot well explore the whole map


How can we efficiently explore an environment **without** using any reward function?

Reward Free Exploration Paradigm

- Reward-free exploration [JKSY20], pure exploration without reward signals
- Two phases algorithm:
 - Exploration Phase: Explore the environment for K episodes
 - No reward signals provided
 - e.g. starting from campus, drive until running out fuel for K times
 - Build estimation of the environment (e.g. road connection in the map)
 - Planning Phase: Given reward signals, plan a near optimal policy
 - No more exploration, only based on estimated environment
 - Can be done multiple times for different reward input
 - e.g. find route to different destinations (UCLA -> LAX, UCLA -> Hollywood, UCLA-> DTLA)

Reward-free exploration: Intuition


- Favoring uncertainty:
 - Explore the environment with more uncertainty!
 - How to measure the uncertainty?
- In tabular setting (i.e. $\mathbb{P}(s' | s, a)$ is explicitly represented by tables)
 - Visiting all states with reasonable probability [JKSY20]
 - Other following up work in tabular setting [MDJKLV20, ZDJ20, KMDJLV21]



We need to explore the areas where we are not sure (by green arrows) starting the initial state (red point)

Function approximation: Mixture of distributions


- Linear Mixture MDPs [JYSW19, AJSWY20, ZHG20]
- Environment $\mathbb{P}(s'|s, a)$ is linear combination of several distributions
- Combining several components of the environment
- $$\mathbb{P}(s'|s, a) = \sum_{i=1}^d \theta_i \mathbb{P}_i(s'|s, a)$$



Simulated car dynamic when turning, with no acceleration

$$\mathbb{P}_1(s'|s, a) \text{ (known)}$$

θ_1



Simulated car dynamic when accelerating

$$\mathbb{P}_2(s'|s, a) \text{ (known)}$$

θ_2



How to estimate the real car's dynamic?

$$\mathbb{P}(s'|s, a) = \theta_1 \mathbb{P}_1(s'|s, a) + \theta_2 \mathbb{P}_2(s'|s, a)$$

A real car with both turning and accelerating

Linear Mixture MDPs: estimating total rewards

- Results from [JYSW19, AJSWY20, ZHG20]:

$$Q_h(s, a) = \underbrace{r_h(s, a)}_{\text{current reward}} + \underbrace{\mathbb{E}[V_{h+1} | s, a]}_{\text{future rewards}}$$

$$\mathbb{E}[V_{h+1} | s, a] = \sum_{i=1}^d \theta_i \mathbb{E}[V_{h+1}(s') | s, a, \mathbb{P}_i] = \langle \boldsymbol{\theta}, \boldsymbol{\psi}_{V_{h+1}}(s, a) \rangle$$

- Plan by Dynamic Programming! $Q_H \rightarrow V_H \rightarrow Q_{H-1} \rightarrow \dots \rightarrow Q_1 \rightarrow V_1$

Uncertainty quantification for Linear Mixture MDPs

- Assume reward functions are known, no uncertainty
- Future uncertainty (Upper Confidence Bound):

$$\left| \hat{\mathbb{E}}[V_{h+1} | s, a] - \mathbb{E}[V_{h+1} | s, a] \right| \leq \beta \sqrt{\boldsymbol{\psi}_{V_{h+1}}^\top(s, a) \Sigma^{-1} \boldsymbol{\psi}_{V_{h+1}}(s, a)}$$

Estimated Expectation

True Expectation

Confidence radius

Covariance matrix

Reward free exploration: our approach (I)

- Encouraging exploration on unknown components
- Exploration driven reward function
 - Intuition: favoring actions lead to more uncertainty


- $r(s, a) \propto \sqrt{\boldsymbol{\psi}_{V_{h+1}}^\top(s, a) \boldsymbol{\Sigma}^{-1} \boldsymbol{\psi}_{V_{h+1}}(s, a)}$

- Issue: no reward, no value function V_{h+1} !

- Solution: $r(s, a) \propto \sqrt{\max_{f: \mathcal{S} \mapsto \mathbb{R}} \boldsymbol{\psi}_f^\top(s, a) \boldsymbol{\Sigma}^{-1} \boldsymbol{\psi}_f(s, a)}$

Exploration Driven reward function: Explained

- $r(s, a) \propto \sqrt{\boldsymbol{\psi}_{V_{h+1}}^\top(s, a) \boldsymbol{\Sigma}^{-1} \boldsymbol{\psi}_{V_{h+1}}(s, a)}$
 - Encouraging exploration with more uncertainty towards a fixed next-step V_{h+1} (e.g. DTLA)
- $r(s, a) \propto \sqrt{\max_{f: \mathcal{S} \mapsto \mathbb{R}} \boldsymbol{\psi}_f^\top(s, a) \boldsymbol{\Sigma}^{-1} \boldsymbol{\psi}_f(s, a)}$
 - Encouraging exploration with more uncertainty globally



Red arrow: exploration towards a fixed goal (DTLA)
Blue arrow: exploration for a global uncertainty

Reward free exploration: our approach (II)

- Using more informative data for regression
- Previous ridge regression in UCRL-VTR [JYSW19, AJSWY20, ZHG20]


$$\mathbb{E}[V(s') | s, a] = \langle \theta^*, \psi_V(s, a) \rangle$$

$$\theta = \arg \min_{\theta} \sum_{(s, a, s')} \left(V(s') - \langle \theta, \psi_V(s, a) \rangle \right)^2 + \lambda \|\theta\|_2^2$$

$V(s')$
Value function
from last step

- Why use V ? Any function can be used to regression!
- Our approach: Uncertainty based regression targets.

Uncertainty based regression targets

- Uncertainty of any function f on fresh sampled data s, a, s' :
 - $\left| \hat{\mathbb{E}}[f(s') | s, a] - \mathbb{E}[f(s') | s, a] \right| \leq \beta \sqrt{\boldsymbol{\psi}_f^\top(s, a) \boldsymbol{\Sigma}^{-1} \boldsymbol{\psi}_f(s, a)}$ 
- Conjecture: larger uncertainty \Leftrightarrow learn more with that target
$$\boldsymbol{\theta} = \arg \min_{\boldsymbol{\theta}} \sum_{(s, a, s')} \left(\mathcal{f}(s') - \langle \boldsymbol{\theta}, \boldsymbol{\psi}_{\mathcal{f}}(s, a) \rangle \right)^2 + \lambda \|\boldsymbol{\theta}\|_2^2$$
- $u = \arg \max_{f: \mathcal{S} \mapsto \mathbb{R}} \left(\boldsymbol{\psi}_{\mathcal{f}}^\top(s, a) \boldsymbol{\Sigma}^{-1} \boldsymbol{\psi}_{\mathcal{f}}(s, a) \right)$ as the target!

Algorithm pipeline: uncertainty based exploration and regression


- Initialize parameter θ , covariance matrix Σ
- Exploration driven reward function in exploration

$$\bullet \quad r(s, a) \propto \sqrt{\max_{f: \mathcal{S} \mapsto \mathbb{R}} \boldsymbol{\psi}_f^\top(s, a) \boldsymbol{\Sigma}^{-1} \boldsymbol{\psi}_f(s, a)}$$

- Sample new data (s, a, s') by maximizing the cumulative reward r
(by UCRL-VTR [JYSW19, AJSWY20, ZHG20])
- Learning the model efficiently using uncertainty

$$u = \arg \max_{f: \mathcal{S} \mapsto \mathbb{R}} \left(\boldsymbol{\psi}_f^\top(s, a) \boldsymbol{\Sigma}^{-1} \boldsymbol{\psi}_f(s, a) \right), \text{ add } u(s') \text{ and } \boldsymbol{\psi}_u(s, a) \text{ into regression}$$

- Update θ, Σ by ridge regression $\theta = \arg \min_{\theta} \sum_{(s, a, s')} (u(s') - \langle \theta, \boldsymbol{\psi}_u(s, a) \rangle)^2 + \lambda \|\theta\|_2^2$



Theoretical Results

Theorem (Sample Complexity)

Let the parameters be properly set, for any $0 < \epsilon < 1$, if $K = \tilde{\mathcal{O}}(H^5d^2\epsilon^{-2})$, with probability at least $1 - \delta$, for any reward function r , the algorithm can provide a policy such that $V_1^*(s; r) - V_1^\pi(s; r) \leq \epsilon$

Take aways

- Collecting $K = \tilde{\mathcal{O}}(H^5d^2\epsilon^{-2})$ is enough to estimate a good enough θ (close to θ^*)
- A well estimated θ is enough to provide a (near) optimal policy for any reward function
- Learning a longer MDP (larger H) is more difficult (why)
 - ★ Longer MDP \Leftrightarrow Larger value function
 - ★ Longer MDP \Leftrightarrow Larger noise in sampling
- Dependency on d and ϵ is standard in linear regression tasks

Comparison to Related Works


Algorithm	Setting	Model Based / Model Free	Sample Complexity
[JKSY20]	Tabular	Model Based	$\tilde{\mathcal{O}}(H^5 S^2 A \epsilon^{-2})$
[KMDJLV21]	Tabular	Model Based	$\tilde{\mathcal{O}}(H^4 S^2 A \epsilon^{-2})$
[MDJKLV20]	Tabular	Model Based	$\tilde{\mathcal{O}}(H^3 S^2 A \epsilon^{-2})$
[ZDJ20]	Tabular	Model Based	$\tilde{\mathcal{O}}(H^2 S^2 A \epsilon^{-2})$
Lower bound [JKSY20]	Tabular	Model Based	$\Omega(H^2 S^2 A \epsilon^{-2})$
[WDYS20]	Linear MDP	Model Free	$\tilde{\mathcal{O}}(H^5 d^3 \epsilon^{-2})$
[ZLKB20]	Linear MDP	Model Free	$\tilde{\mathcal{O}}(H^5 d^3 \epsilon^{-2})$
Lower bound [WDYS20]	Linear MDP	Model Free	Not Applicable
UCRL-RFE (ours)	Linear Mixture MDP	Model Based	$\tilde{\mathcal{O}}(H^5 d^2 \epsilon^{-2})$
UCRL-RFE+ (ours, improved)	Linear Mixture MDP	Model Based	$\tilde{\mathcal{O}}(H^4 d(H + d) \epsilon^{-2})$
Lower bound (ours)	Linear Mixture MDP	Model Based	$\Omega(H^2 d \epsilon^{-2})$
Lower bound (Improved) [CHYW21]	Linear Mixture MDP	Model Based	$\Omega(H^2 d^2 \epsilon^{-2})$

Sample complexity of the reward free exploration algorithms, the time-inhomogeneous results are translated to time-homogeneous results by removing

an H factor

Experiment Results

- $d = 3, \mathcal{S} = 10, \mathcal{A} = 5, H = 10, K = 100$
- Random Exploration:
 - Random exploration policy
 - Random regression target
- Oracle (UCRL-VTR) [JYSW19, AJSWY20, ZHG20]:
 - Know the reward function during exploration
- Reward Free Exploration (Ours):
 - Reward free exploration is significantly better than random exploration!



Further extensions

- Explore to non-linear function approximation
 - Once we know the uncertainty, we can use that to guide exploration
- Improve the theoretical bounds (Current upper bound: $\mathcal{O}(H^5d^2\epsilon^{-2})$)
 - Lower bound is improved to $\Omega(H^2d^2\epsilon^{-2})$ from $\Omega(H^2d\epsilon^{-2})$ [CHYW21]
 - The dependency on H still have a large gap
- Empirical results can be applied to modern model-based algorithms.

Summary

Quantify uncertainty of RL with function approximation can

- Guide the exploration in RL
- Improve the efficiency of learning the parameters

The uncertainty of function approximated RL can be quantified by ...

- Precisely controlled with linear function approximation
- Easy to immigrate to other models (e.g. neural networks [ZZLQ20])

My future plan

- Exploration: Pure exploration without reward signals
- Model Misspecification Issue: Control the approximation error in the model
- Representation Learning: Select good representation to improve performance

- Partially observed RL and non-Markovian RL: Missing information in current observation
- Fairness in RL: Make fair decision when uncertainty exists
- Deep RL: quantify the uncertainty in neural networks used in RL
- Practical algorithms using modern neural networks, on modern tasks..

Thank you!